

Question 1: Which of the following is the correct general structure of a typical process?

- A) Initialization, Execution, Termination
- B) Execution, Termination, Initialization
- C) Termination, Initialization, Execution
- D) Initialization, Termination, Execution

Answer: A) Initialization, Execution, Termination

Question 2: A race condition in computer science occurs when:

- A) Two or more processes accidentally cooperate to complete a task.
- B) Two or more processes contend for shared resources and the final outcome depends on their relative timing.
- C) Two or more processes work in parallel without any synchronization.
- D) Two or more processes try to access the same resource sequentially.

Answer: B) Two or more processes contend for shared resources and the final outcome depends on their relative timing.

Question 3: The Critical-Section Problem involves:

- A) Allowing multiple processes to access a critical section simultaneously.
- B) Ensuring that only one process can access a critical section at a time.
- C) Letting processes skip the critical section for faster execution.
- D) Preventing processes from entering the initialization phase.

Answer: B) Ensuring that only one process can access a critical section at a time.

Question 4: Which of the following is a solution to the Critical-Section Problem proposed by Peterson?

- A) Mutex Locks
- B) Semaphores
- C) Monitors
- D) Two-Process Solution

Answer: D) Two-Process Solution (Peterson's Solution)

Question 5: Synchronization hardware such as Test-And-Set and Compare-And-Swap instructions are used for:

- A) Achieving deadlock in concurrent processes.
- B) Speeding up critical section execution.
- C) Solving the Critical-Section Problem without mutual exclusion.
- D) Implementing atomic operations to manage access to shared resources.

Answer: D) Implementing atomic operations to manage access to shared resources.

Question 6: A Mutex Lock is:

- A) A type of data structure used for storing multiple processes.
- B) A synchronization primitive that allows multiple processes to access a critical section simultaneously.
- C) A synchronization mechanism that ensures only one process can access a critical section at a time.
- D) A hardware component used for inter-process communication.

Answer: C) A synchronization mechanism that ensures only one process can access a critical section at a time.

Question 7: Semaphores can be used to:

- A) Allow processes to run without synchronization.
- B) Induce race conditions in concurrent systems.
- C) Implement mutual exclusion and coordinate concurrent processes.
- D) Replace the need for critical sections in programs.

Answer: C) Implement mutual exclusion and coordinate concurrent processes.

Question 8: Which of the following is an example of a classic synchronization problem?

- A) Sorting a list of numbers.
- B) Performing matrix multiplication.
- C) Dining Philosophers Problem.
- D) File I/O operations.

Answer: C) Dining Philosophers Problem

Question 9: A Monitor is a synchronization construct that:

- A) Prevents context switching between processes.
- B) Enables processes to run in parallel without synchronization.
- C) Provides a high-speed communication channel between processes.
- D) Encapsulates shared data and procedures to control their access in a synchronized manner.

Answer: D) Encapsulates shared data and procedures to control their access in a synchronized manner.

Question 10: What is the primary goal of process synchronization in operating systems?

- A) To increase the number of processes running concurrently.
- B) To ensure that all processes complete their execution at the same time.
- C) To coordinate the execution of multiple processes and avoid conflicts.
- D) To prioritize certain processes over others.

Answer: C) To coordinate the execution of multiple processes and avoid conflicts.

Question 11: Which of the following is a potential issue that can arise from improper process synchronization?

- A) Faster execution of critical sections.
- B) Enhanced parallelism among processes.
- C) Deadlocks and race conditions.
- D) Decreased utilization of CPU.

Answer: C) Deadlocks and race conditions.

Question 12: A race condition can occur when:

- A) Processes follow a predetermined order of execution.
- B) Processes synchronize their execution using semaphores.
- C) Processes have inconsistent access to shared resources.
- D) Processes use hardware synchronization mechanisms.

Answer: C) Processes have inconsistent access to shared resources.

Question 13: Peterson's solution for the Critical-Section Problem is best suited for:

- A) Systems with only two processes.
- B) Systems with a small number of processes.
- C) Systems with a large number of processes.
- D) Systems with processes of varying priorities.

Answer: A) Systems with only two processes.

Question 14: The term starvation in synchronization refers to:

- A) A process not being allowed to execute its critical section.
- B) A process completing its execution too quickly.
- C) A process being stuck in a deadlock.
- D) A process experiencing excessive context switching.

Answer: A) A process not being allowed to execute its critical section.

Question 15: In synchronization, a mutex is:

- A) A variable used to count the number of processes.
- B) A hardware device used to prevent deadlocks.
- C) A synchronization primitive that ensures exclusive access to a shared resource.
- D) A type of process that operates without synchronization.

Answer: C) A synchronization primitive that ensures exclusive access to a shared resource.

Question 16: Which of the following statements about semaphores is true?

- A) Semaphores are used only for signaling purposes.
- B) Binary semaphores can have values greater than 1.
- C) Semaphores can be used for both signaling and mutual exclusion.
- D) Semaphores are primarily used for thread synchronization.

Answer: C) Semaphores can be used for both signaling and mutual exclusion.

Question 17: In a Producer-Consumer synchronization problem, the main challenge is to:

- A) Prevent deadlocks between producers and consumers.
- B) Ensure that producers always produce faster than consumers consume.
- C) Coordinate access to shared resources between multiple producers and consumers.
- D) Limit the number of producers and consumers in the system.

Answer: C) Coordinate access to shared resources between multiple producers and consumers.

Question 18: The Readers-Writers problem involves:

- A) Ensuring that multiple readers can access a shared resource simultaneously.
- B) Allowing only one writer to access a shared resource at a time.
- C) Preventing readers and writers from accessing the shared resource.
- D) Allowing readers to read and writers to write concurrently without synchronization.

Answer: D) Allowing readers to read and writers to write concurrently without synchronization.

Question 19: What is the purpose of a monitor in synchronization?

- A) To provide a graphical representation of process interactions.
- B) To prevent processes from entering the critical section.
- C) To encapsulate shared data and provide synchronized access to it.
- D) To execute processes in a concurrent manner.

Answer: C) To encapsulate shared data and provide synchronized access to it.

Question 20: The Two-Phase Locking protocol is primarily used to address:

- A) Deadlock prevention.
- B) Starvation in synchronization.
- C) Mutual exclusion.
- D) Deadlock detection.

Answer: A) Deadlock prevention.

Question 21: In synchronization, a condition variable is used to:

- A) Store the value of a synchronization primitive.
- B) Indicate the status of the CPU.
- C) Synchronize access to shared resources.
- D) Enable efficient context switching.

Answer: C) Synchronize access to shared resources.

Question 22: Which of the following is an example of a classic synchronization problem involving multiple processes and dining philosophers?

- A) Readers-Writers Problem
- B) Producer-Consumer Problem
- C) Dining Philosophers Problem
- D) Sleeping Barber Problem

Answer: C) Dining Philosophers Problem

Question 23: A deadlock occurs when:

- A) A process finishes its execution.
- B) A process enters its critical section.
- C) Processes are waiting for a resource held by another process.
- D) Processes are executing in parallel.

Answer: C) Processes are waiting for a resource held by another process.

Question 24: The term atomic operation in synchronization refers to:

- A) An operation that occurs only once in the lifetime of a process.
- B) An operation that is executed by multiple processes simultaneously.
- C) An operation that is indivisible and uninterruptible.
- D) An operation that is performed by the CPU in a single clock cycle.

Answer: C) An operation that is indivisible and uninterruptible.

Question 25: What is the primary purpose of the test-and-set instruction in synchronization?

- A) To test the performance of a CPU.
- B) To set a flag indicating the status of a process.

- C) To perform bitwise operations on data.
- D) To implement atomic operations and manage access to shared resources.

Answer: D) To implement atomic operations and manage access to shared resources.

Question 26: A semaphore with a value of 0 indicates:

- A) That a resource is available for use.
- B) That the semaphore is in a waiting state.
- C) That the semaphore is in an undefined state.
- D) That a process must wait for a resource to be released.

Answer: D) That a process must wait for a resource to be released.

Question 27: What is the primary purpose of the Compare-And-Swap (CAS) instruction in synchronization?

- A) To compare two variables and perform a swap operation if they are equal.
- B) To perform arithmetic operations on data.
- C) To swap the contents of two memory locations.
- D) To perform context switching between processes.

Answer: A) To compare two variables and perform a swap operation if they are equal.

Question 28: A mutex lock ensures:

- A) That multiple processes can access a critical section concurrently.
- B) That processes are locked out of the system.
- C) That only one process can access a critical section at a time.
- D) That processes execute in a random order.

Answer: C) That only one process can access a critical section at a time.

Question 29: In the context of synchronization, what is a priority inversion?

- A) A high-priority process getting locked out by a lower-priority process.
- B) A low-priority process gaining priority over a high-priority process.
- C) A process getting stuck in a deadlock due to priority conflicts.
- D) A process finishing its execution before others due to priority adjustments.

Answer: A) A high-priority process getting locked out by a lower-priority process.

Question 30: The Bakery Algorithm is used for:

- A) Implementing mutual exclusion between two processes.
- B) Preventing priority inversion in synchronization.
- C) Solving the Dining Philosophers Problem.
- D) Coordinating access to a critical section among multiple processes.

Answer: D) Coordinating access to a critical section among multiple processes.

Question 31: In a Monitor, what is the purpose of a condition variable?

- A) To store the monitor's internal state.
- B) To define the entry point of the monitor.
- C) To signal other processes to exit the monitor.
- D) To allow processes to wait for a certain condition to be satisfied.

Answer: D) To allow processes to wait for a certain condition to be satisfied.

Question 32: The Three-State Process Model includes which of the following states?

- A) Running, Sleeping, Blocking
- B) Running, Waiting, Terminated
- C) Ready, Running, Blocked
- D) Ready, Sleeping, Blocked

Answer: C) Ready, Running, Blocked

Question 33: In a system with multiple processes, the term context switch refers to:

- A) Switching between different processors in a multi-core CPU.
- B) The process of changing the execution context from one process to another.
- C) The act of rebooting the system.
- D) Switching between different synchronization mechanisms.

Answer: B) The process of changing the execution context from one process to another.

Question 34: What is the primary purpose of the turnstile synchronization mechanism?

- A) To prevent processes from executing their critical sections.
- B) To manage the flow of processes in a round-robin fashion.
- C) To ensure that processes exit their critical sections in a timely manner.
- D) To enforce a certain order of execution among processes.

Answer: C) To ensure that processes exit their critical sections in a timely manner.

Question 35: Which of the following statements is true regarding synchronization primitives?

- A) They are only used for debugging purposes.
- B) They are low-level mechanisms used to coordinate processes.
- C) They are only used in single-processor systems.
- D) They are primarily used for improving CPU performance.

Answer: B) They are low-level mechanisms used to coordinate processes.

Question 36: What is the primary difference between a mutex and a semaphore?

- A) A mutex allows multiple processes to access a critical section at once, while a semaphore doesn't.
- B) A mutex can be initialized with a value greater than 1, while a semaphore cannot.
- C) A mutex is binary, allowing only one process at a time, while a semaphore can have multiple values.
- D) A mutex is a hardware synchronization mechanism, while a semaphore is a software construct.

Answer: C) A mutex is binary, allowing only one process at a time, while a semaphore can have multiple values.

Question 37: The Two-Phase Commit (2PC) protocol is used for:

- A) Coordinating the access to a critical section among multiple processes.
- B) Solving the Dining Philosophers Problem.
- C) Ensuring that multiple processes commit their transactions atomically.
- D) Preventing priority inversion in synchronization.

Answer: C) Ensuring that multiple processes commit their transactions atomically.

Question 38: In a Peterson's Solution for the Critical-Section Problem, what is the purpose of the turn variable?

- A) To indicate which process should run next.
- B) To determine the order of execution among multiple processes.
- C) To prevent deadlock situations.
- D) To signal when a process enters the critical section.

Answer: B) To determine the order of execution among multiple processes.

Question 39: Which of the following is NOT a common approach to solving synchronization problems?

- A) Using mutex locks
- B) Utilizing semaphores
- C) Ignoring synchronization altogether
- D) Employing condition variables

Answer: C) Ignoring synchronization altogether

Question 40: In a Monitor, why are procedures that access shared data usually declared as synchronized?

- A) To ensure that only one process can access a procedure at a time.
- B) To improve the execution speed of the procedures.
- C) To allow the procedures to run concurrently without synchronization.
- D) To allow the procedures to be used as entry points to the monitor.

Answer: A) To ensure that only one process can access a procedure at a time.

Question 1: What is the primary goal of CPU scheduling in operating systems?

- A) To maximize the throughput of the CPU.
- B) To minimize the execution time of individual processes.
- C) To ensure that all processes get equal execution time.
- D) To efficiently allocate CPU time to multiple processes.

Answer: D) To efficiently allocate CPU time to multiple processes.

Question 2: Which of the following is a common scheduling criterion used to evaluate scheduling algorithms?

- A) Maximizing CPU utilization
- B) Minimizing waiting time
- C) Minimizing turnaround time
- D) All of the above

Answer: D) All of the above

Question 3: In the First-Come, First-Served (FCFS) scheduling algorithm, processes are executed in the order they arrive. What is a major drawback of this algorithm?

- A) It doesn't guarantee fairness among processes.
- B) It can lead to starvation.
- C) It doesn't maximize CPU utilization.
- D) It doesn't support multi-core processors.

Answer: B) It can lead to starvation.

Question 4: Shortest Job First (SJF) scheduling aims to:

- A) Give priority to processes with the shortest burst time.
- B) Ensure that processes with the longest burst time get executed first.
- C) Schedule processes based on their arrival time.
- D) Alternate between processes in a round-robin manner.

Answer: A) Give priority to processes with the shortest burst time.

Question 5: What is the main difference between Shortest Job First (SJF) and Shortest Remaining Time First (SRTF) scheduling?

- A) SJF is non-preemptive, while SRTF is preemptive.
- B) SJF is preemptive, while SRTF is non-preemptive.
- C) Both SJF and SRTF are non-preemptive.
- D) Both SJF and SRTF are preemptive.

Answer: A) SJF is non-preemptive, while SRTF is preemptive.

Question 6: The Priority Scheduling algorithm assigns priorities to processes based on:

- A) Their arrival time.
- B) Their execution time.
- C) User-defined priority values.
- D) The order in which they appear in the queue.

Answer: C) User-defined priority values.

Question 7: Which scheduling algorithm is designed to address the Convoy Effect by allowing short processes to complete quickly?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: B) SJF

Question 8: In Round Robin (RR) scheduling, each process is assigned a fixed time slice called a quantum. What happens if a process doesn't complete within its time quantum?

- A) It gets terminated immediately.
- B) It is moved to the end of the queue.
- C) It continues executing without interruption.
- D) It is given additional time slices to complete.

Answer: B) It is moved to the end of the queue.

Question 9: What is the primary advantage of Multilevel Queue Scheduling over other scheduling algorithms?

- A) It guarantees that every process gets an equal share of CPU time.
- B) It allows processes to be assigned to different priority levels.
- C) It eliminates the need for context switching.



D) It supports only single-core processors.

Answer: B) It allows processes to be assigned to different priority levels.

Question 10: In Multilevel Feedback Queue Scheduling, how are processes promoted or demoted between queues?

- A) Based on their execution time
- B) Based on their priority levels
- C) Based on their arrival time
- D) Based on the number of CPUs available

Answer: A) Based on their execution time

Question 11: What is the purpose of Thread Scheduling in a multi-threaded environment?

- A) To allocate memory to threads.
- B) To allocate CPU time to threads.
- C) To manage synchronization between threads.
- D) To prioritize thread execution based on their memory requirements.

Answer: B) To allocate CPU time to threads.

Question 12: The Convoy Effect refers to:

- A) A situation where a large number of processes are waiting to be executed.
- B) The process of moving a process to the back of the queue after its time quantum expires.
- C) Processes with similar burst times arriving simultaneously.
- D) A high-priority process blocking the execution of lower-priority processes.

Answer: A) A situation where a large number of processes are waiting to be executed.

Question 13: Which scheduling algorithm is considered optimal in terms of minimizing average waiting time?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: B) SJF

Question 14: In Priority Scheduling, what can be a potential issue if a high-priority process continues to arrive frequently?

- A) Starvation of low-priority processes
- B) Inefficient utilization of the CPU
- C) Excessive context switching
- D) Priority inversion

Answer: A) Starvation of low-priority processes

Question 15: What is the significance of the time quantum in the Round Robin scheduling algorithm?

- A) It determines the priority of a process.
- B) It specifies the time a process must wait before it can execute.
- C) It defines the maximum burst time for a process.

D) It limits the maximum amount of time a process can execute before being moved to the end of the queue.

Answer: D) It limits the maximum amount of time a process can execute before being moved to the end of the queue.

Question 16: Which scheduling algorithm is suitable for systems that prioritize interactive processes and ensure that no process waits too long?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: D) Round Robin

Question 17: In Multilevel Feedback Queue Scheduling, what determines the priority of a process in different queues?

- A) The arrival time of the process
- B) The execution time of the process
- C) The type of the process (CPU-bound or I/O-bound)
- D) The number of threads the process uses

Answer: B) The execution time of the process

Question 18: What is the primary advantage of using Thread Scheduling in a multi-threaded application?

- A) It reduces the memory usage of the application.
- B) It allows threads to run in parallel on separate processors.
- C) It improves the overall performance of the application.
- D) It eliminates the need for synchronization between threads.

Answer: B) It allows threads to run in parallel on separate processors.

Question 19: Which scheduling algorithm is well-suited for real-time systems where tasks have strict deadlines to meet?

- A) FCFS
- B) SJF
- C) Priority S

cheduling

- D) Round Robin

Answer: C) Priority Scheduling

Question 20: The aging mechanism in Priority Scheduling aims to:

- A) Increase the priority of low-priority processes.
- B) Gradually reduce the priority of processes over time.
- C) Prevent high-priority processes from executing too frequently.
- D) Promote short processes to higher-priority levels.

Answer: A) Increase the priority of low-priority processes.

Question 21:What is the primary drawback of the First-Come, First-Served (FCFS) scheduling algorithm?

- A) It can result in low CPU utilization.
- B) It does not support multi-core processors.
- C) It favors short processes over long processes.
- D) It can lead to the convoy effect.

Answer: A) It can result in low CPU utilization.

Question 22:In the context of CPU scheduling, what does the term context switch refer to?

- A) Moving a process from the waiting queue to the ready queue.
- B) Changing the priority of a process.
- C) Saving the current state of a running process and loading the state of a new process.
- D) Allocating CPU time to a process.

Answer: C) Saving the current state of a running process and loading the state of a new process.

Question 23:The scheduling algorithm that allows a process to run for a fixed time quantum and then moves it to the back of the queue is:

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: D) Round Robin

Question 24:Which scheduling algorithm is suitable for time-sharing systems, where multiple users access the system concurrently?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: D) Round Robin

Question 25:In Multilevel Queue Scheduling, how are processes categorized into different queues?

- A) Based on their arrival time
- B) Based on their execution time
- C) Based on their memory requirements
- D) Based on their priority or type of task

Answer: D) Based on their priority or type of task

Question 26:Which scheduling algorithm has the potential to suffer from priority inversion?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: C) Priority Scheduling

Question 27: In Multilevel Feedback Queue Scheduling, which queue has the highest priority?

- A) The queue with the shortest burst time processes
- B) The queue with the longest burst time processes
- C) The queue with the highest priority processes
- D) The queue with the lowest priority processes

Answer: C) The queue with the highest priority processes

Question 28: What is the purpose of using a time quantum in the Round Robin scheduling algorithm?

- A) To determine the order of execution among processes.
- B) To ensure fairness among processes.
- C) To minimize the context switch overhead.
- D) To ensure that each process gets an equal share of the CPU time.

Answer: D) To ensure that each process gets an equal share of the CPU time.

Question 29: Which scheduling algorithm is most appropriate for a system where all processes have the same priority?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: A) FCFS

Question 30: What is the primary advantage of Shortest Remaining Time First (SRTF) scheduling over Shortest Job First (SJF) scheduling?

- A) SRTF provides better throughput.
- B) SRTF eliminates the need for context switching.
- C) SRTF reduces the number of context switches.
- D) SRTF is more responsive in interactive environments.

Answer: D) SRTF is more responsive in interactive environments.

Question 31: The scheduling algorithm that allows processes to execute in a circular order and provides fairness among processes is:

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: D) Round Robin

Question 32: What is the primary limitation of using Round Robin scheduling with a large time quantum?

- A) It leads to excessive context switching.
- B) It results in low CPU utilization.
- C) It can cause high waiting times for processes.
- D) It makes the scheduling algorithm too complex.

Answer: B) It results in low CPU utilization.

Question 33: In a Multilevel Feedback Queue scheduling algorithm, a process that uses the CPU extensively will:

- A) Be demoted to a lower-priority queue.
- B) Be promoted to a higher-priority queue.
- C) Continue executing in the same queue.
- D) Be moved to the end of the queue.

Answer: A) Be demoted to a lower-priority queue.

Question 34: What is the key advantage of Thread Scheduling over Process Scheduling?

- A) Thread scheduling reduces memory usage.
- B) Thread scheduling eliminates the need for context switches.
- C) Thread scheduling allows better utilization of multi-core processors.
- D) Thread scheduling improves the responsiveness of the system.

Answer: C) Thread scheduling allows better utilization of multi-core processors.

Question 35: Which scheduling algorithm guarantees the lowest average waiting time for a set of processes?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: B) SJF

Question 36: In Priority Scheduling, what happens if two processes have the same priority?

- A) The process with the shortest burst time is given priority.
- B) The process that arrives first is given priority.
- C) The process that uses the CPU the most is given priority.
- D) The process with the longest burst time is given priority.

Answer: B) The process that arrives first is given priority.

Question 37: Which scheduling algorithm ensures that each process gets an equal share of the CPU time?

- A) FCFS
- B) SJF
- C) Priority Scheduling
- D) Round Robin

Answer: D) Round Robin

Question 38: The Multilevel Feedback Queue Scheduling algorithm uses different queues for processes based on:

- A) The number of threads a process has.
- B) The arrival time of processes.
- C) The burst time of processes.
- D) The execution history of processes.

Answer: D) The execution history of processes.

Question 39:What is the primary benefit of Thread Scheduling in a multi-threaded environment?

- A) It eliminates the need for synchronization.
- B) It reduces the overhead of creating new processes.
- C) It allows threads to run concurrently on separate processors.
- D) It enforces strict control over CPU utilization.

Answer: C) It allows threads to run concurrently on separate processors.

Question 40:The Convoy Effect occurs when:

- A) A high-priority process blocks the execution of lower-priority processes.
- B) Processes with similar burst times arrive simultaneously.
- C) A long process gets executed before short processes.
- D) Several processes wait for a resource held by another process.

Answer: C) A long process gets executed before short processes.